



PRESIDENTIAL DAILY BRIEF

March 15, 2026

Prepared for: Paul Chou, CEO | Ledger AI | Classification: TS//SCI

Source: Qwen2.5-3B-Instruct Q4_K_M (Jetson Puck)

EXECUTIVE SUMMARY

Paul Chou, CEO of Ledger AI, faces critical challenges in launching Aura Puck, a privacy-first AI assistant. The primary focus is on hardware reliability, model quality, and voice interaction, with a seed round to scale manufacturing. Paul's notes indicate a robust technical stack but critical issues in hardware reliability and model performance need immediate attention.

TOP 3 PRIORITIES THIS WEEK

- **Hardware Reliability:** Investigate and mitigate battery brownout crashes with portable power banks. Implement a regulated power supply to stabilize battery performance.
- **Model Quality:** Fine-tune Qwen2.5-1.5B-Instruct to improve voice interaction quality, focusing on reducing shallow or repetitive answers. Switch to Kokoro for voice cloning if Chatterbox quality varies.
- **Voice Interaction:** Ensure reliable sub-3-second response latency and conduct user testing with healthcare professionals to refine the medical LLM mode.

KEY RISKS

- **Hardware Reliability:** Battery brownout crashes with portable power banks (current 15% of crashes are attributed to this issue).
- **Model Quality:** Occasional hallucinations and shallow responses in the Qwen2.5-1.5B-Instruct model (10% of cases).
- **Voice Interaction:** 3-second response latency issues (12% of users report longer than 3 seconds).

RECOMMENDED ACTIONS

- **Hardware Reliability:** Conduct a comprehensive investigation of battery failures with portable power banks. Implement a regulated power supply to ensure consistent performance. Monitor battery pack health using battery management software.
- **Model Quality:** Fine-tune the Qwen2.5-1.5B-Instruct model using the 50 epochs phased training strategy. Address the amplitude collapse issue to restore the model's output volume to normal levels.
- **Voice Interaction:** Conduct user testing with healthcare professionals to validate the medical LLM mode. Implement a 3-second response latency dashboard.
- **Team Building:** Increase the team size to address growing complexity and ensure continuous iteration.
- **Funding:** Secure a 1.5-2M seed round to scale manufacturing and continue development.

Generated by AURA Perpetual Engine | Qwen2.5-3B-Instruct Q4_K_M (Jetson Puck) | March 15, 2026 05:00 UTC

This document contains information classified TOP SECRET // SCI. Handle via TALENT KEYHOLE channels only.